

```
In [4]: import tensorflow as tf
import tensorflow.keras import models, layers
import matplotlib.pyplot as plt
from IPython.display import HTML

In [5]: BATCH_SIZE = 32
IMAGE_SIZE = 256
CHANNELS=3
EPOCHS=50

In [9]: dataset = tf.keras.preprocessing.image_dataset_from_directory(
    "downloads/Plantvillage",
    seed=123,
    shuffle=True,
    image_size=(IMAGE_SIZE, IMAGE_SIZE),
    batch_size=BATCH_SIZE
)

Found 2152 files belonging to 4 classes.

In [11]: class_names = dataset.class_names
class_names

Out[11]: ['_ipynb_checkpoints',
'Potato__Early_blight',
'Potato__Late_blight',
'Potato__healthy']

In [12]: for image_batch, labels_batch in dataset.take(1):
    print(image_batch.shape)
    print(labels_batch.numpy())

(32, 256, 256, 3)
[2 2 2 1 1 1 1 1 2 2 2 1 2 1 2 2 1 2 1 2 1 1 2 2 3 1 1]

In [13]: #Visualize some of the images from our dataset

In [14]: plt.figure(figsize=(10, 10))
for image_batch, labels_batch in dataset.take(1):
    for i in range(12):
        ax = plt.subplot(3, 4, i + 1)
        plt.imshow(image_batch[i].numpy().astype("uint8"))
        plt.title(class_names[labels_batch[i]])
        plt.axis("off")

Potato__Early_blight
Potato__Early_blight
Potato__Early_blight
Potato__Late_blight

Potato__Early_blight
Potato__Early_blight
Potato__Late_blight
Potato__Early_blight

Potato__Late_blight
Potato__Early_blight
Potato__Early_blight
Potato__Early_blight

In [15]: len(dataset)

Out[15]: 68

In [16]: train_size = 0.8
len(dataset)*train_size

Out[16]: 54.400000000000006

In [17]: train_ds = dataset.take(54)
len(train_ds)

Out[17]: 54

In [18]: test_ds = dataset.skip(54)
len(test_ds)

Out[18]: 14

In [19]: val_size=0.1
len(dataset)*val_size

Out[19]: 6.800000000000001

In [20]: val_ds = test_ds.take(6)
len(val_ds)

Out[20]: 6

In [21]: test_ds = test_ds.skip(6)
len(test_ds)

Out[21]: 8

In [22]: def get_dataset_partitions_tf(ds, train_split=0.8, val_split=0.1, test_split=0.1, shuffle=True, shuffle_size=10000):
    assert (train_split + test_split + val_split) == 1

    ds_size = len(ds)

    if shuffle:
        ds = ds.shuffle(shuffle_size, seed=12)

    train_size = int(train_split * ds_size)
    val_size = int(val_split * ds_size)

    train_ds = ds.take(train_size)
    val_ds = ds.skip(train_size).take(val_size)
    test_ds = ds.skip(train_size).skip(val_size)

    return train_ds, val_ds, test_ds

In [23]: train_ds, val_ds, test_ds = get_dataset_partitions_tf(dataset)

In [24]: len(train_ds)

Out[24]: 54

In [25]: len(val_ds)

Out[25]: 6

In [26]: len(test_ds)

Out[26]: 8

In [27]: #Cache, Shuffle, and Prefetch the Dataset

In [28]: train_ds = train_ds.cache().shuffle(1000).prefetch(buffer_size=tf.data.AUTOTUNE)
val_ds = val_ds.cache().shuffle(1000).prefetch(buffer_size=tf.data.AUTOTUNE)
test_ds = test_ds.cache().shuffle(1000).prefetch(buffer_size=tf.data.AUTOTUNE)

In [29]: resize_and_rescale = tf.keras.Sequential([
    layers.experimental.preprocessing.Resizing(IMAGE_SIZE, IMAGE_SIZE),
    layers.experimental.preprocessing.Rescaling(1./255),
])

In [30]: #Data Augmentation
#Data Augmentation is needed when we have less data, this boosts the accuracy of our model by augmenting the data.

In [31]: data_augmentation = tf.keras.Sequential([
    layers.experimental.preprocessing.RandomFlip("horizontal_and_vertical"),
    layers.experimental.preprocessing.RandomRotation(0.2),
])

In [32]: #Applying Data Augmentation to Train Dataset

In [33]: train_ds = train_ds.map(
    lambda x, y: (data_augmentation(x, training=True), y)
).prefetch(buffer_size=tf.data.AUTOTUNE)

WARNING:tensorflow:From C:\Users\COMP\anaconda3\lib\site-packages\tensorflow\python\autograph\pyct\static_analysis\liveness.py:83: Analyzer.lamba_check (from tensorflow.python.autograph.pyct.static_analysis.liveness) is deprecated and will be removed after 2023-09-23.
Instructions for updating:
Lambda fuctions will be no more assumed to be used in the statement where they are used, or at least in the same block. https://github.com/tensorflow/tensorflow/issues/56089
WARNING:tensorflow:Using a while_loop for converting RngReadAndSkip cause there is no registered converter for this op.
WARNING:tensorflow:Using a while_loop for converting Bitcast cause there is no registered converter for this op.
WARNING:tensorflow:Using a while_loop for converting Bitcast cause there is no registered converter for this op.
WARNING:tensorflow:Using a while_loop for converting StatelessRandomUniformV2 cause there is no registered converter for this op.
WARNING:tensorflow:Using a while_loop for converting ImageProjectiveTransformV3 cause there is no registered converter for this op.
WARNING:tensorflow:Using a while_loop for converting RngReadAndSkip cause there is no registered converter for this op.
WARNING:tensorflow:Using a while_loop for converting Bitcast cause there is no registered converter for this op.
WARNING:tensorflow:Using a while_loop for converting Bitcast cause there is no registered converter for this op.
WARNING:tensorflow:Using a while_loop for converting StatelessRandomUniformV2 cause there is no registered converter for this op.
WARNING:tensorflow:Using a while_loop for converting ImageProjectiveTransformV3 cause there is no registered converter for this op.

In [34]: input_shape = (BATCH_SIZE, IMAGE_SIZE, IMAGE_SIZE, CHANNELS)
n_classes = 3

model = models.Sequential([
    resize_and_rescale,
    layers.Conv2D(32, kernel_size = (3,3), activation='relu', input_shape=input_shape),
    layers.MaxPooling2D((2, 2)),
    layers.Conv2D(64, kernel_size = (3,3), activation='relu'),
    layers.MaxPooling2D((2, 2)),
    layers.Conv2D(64, kernel_size = (3,3), activation='relu'),
    layers.MaxPooling2D((2, 2)),
    layers.Conv2D(64, (3, 3), activation='relu'),
    layers.MaxPooling2D((2, 2)),
    layers.Conv2D(64, (3, 3), activation='relu'),
    layers.MaxPooling2D((2, 2)),
    layers.Flatten(),
    layers.Dense(64, activation='relu'),
    layers.Dense(n_classes, activation='softmax'),
])

model.build(input_shape=input_shape)

In [35]: model.summary()

Model: "sequential_2"

Layer (type) Output Shape Param #
-----
sequential (Sequential) (32, 256, 256, 3) 0
conv2d (Conv2D) (32, 254, 254, 32) 896
max_pooling2d (MaxPooling2D) (32, 127, 127, 32) 0
conv2d_1 (Conv2D) (32, 125, 125, 64) 18496
max_pooling2d_1 (MaxPooling2D) (32, 62, 62, 64) 0
conv2d_2 (Conv2D) (32, 60, 60, 64) 36928
max_pooling2d_2 (MaxPooling2D) (32, 30, 30, 64) 0
conv2d_3 (Conv2D) (32, 28, 28, 64) 36928
max_pooling2d_3 (MaxPooling2D) (32, 14, 14, 64) 0
conv2d_4 (Conv2D) (32, 12, 12, 64) 36928
max_pooling2d_4 (MaxPooling2D) (32, 6, 6, 64) 0
conv2d_5 (Conv2D) (32, 4, 4, 64) 36928
max_pooling2d_5 (MaxPooling2D) (32, 2, 2, 64) 0
flatten (Flatten) (32, 256) 0
dense (Dense) (32, 64) 16448
dense_1 (Dense) (32, 3) 195

Total params: 183,747
Trainable params: 183,747
Non-trainable params: 0

In [36]: #Compiling the Model

In [37]: model.compile(
    optimizer='adam',
    loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=False),
    metrics=['accuracy']
)

In [47]: import numpy as np
for images_batch, labels_batch in test_ds.take(1):

    first_image = images_batch[0].numpy().astype('uint8')
    first_label = labels_batch[0].numpy()

    print("first image to predict")
    plt.imshow(first_image)
    print("actual label:", class_names[first_label])

    batch_prediction = model.predict(images_batch)
    print("predicted label:", class_names[np.argmax(batch_prediction[0])])

first image to predict
actual label: Potato__Late_blight
1/1 [=====] - 2s 2s/step
predicted label: Potato__Early_blight

0
50
100
150
200
250
0 50 100 150 200 250

In [48]: def predict(model, img):
    img_array = tf.keras.preprocessing.image.img_to_array(images[1].numpy())
    img_array = tf.expand_dims(img_array, 0)

    predictions = model.predict(img_array)

    predicted_class = class_names[np.argmax(predictions[0])]
    confidence = round(100 * (np.max(predictions[0])), 2)
    return predicted_class, confidence

In [49]: plt.figure(figsize=(15, 15))
for images, labels in test_ds.take(1):
    for i in range(9):
        ax = plt.subplot(3, 3, i + 1)
        plt.imshow(images[1].numpy().astype("uint8"))

        predicted_class, confidence = predict(model, images[1].numpy())
        actual_class = class_names[labels[1]]

        plt.title(f"Actual: {actual_class},\n Predicted: {predicted_class},\n Confidence: {confidence}%")

        plt.axis("off")

1/1 [=====] - 0s 343ms/step
1/1 [=====] - 0s 70ms/step
1/1 [=====] - 0s 72ms/step
1/1 [=====] - 0s 87ms/step
1/1 [=====] - 0s 80ms/step
1/1 [=====] - 0s 95ms/step
1/1 [=====] - 0s 86ms/step
1/1 [=====] - 0s 83ms/step
1/1 [=====] - 0s 86ms/step

Actual: Potato__Late_blight,
Predicted: Potato__Early_blight.
Confidence: 33.95%

Actual: Potato__Early_blight,
Predicted: Potato__Early_blight.
Confidence: 34.05%

Actual: Potato__Early_blight,
Predicted: Potato__Early_blight.
Confidence: 33.95%

Actual: Potato__Early_blight,
Predicted: Potato__Early_blight.
Confidence: 34.08%

Actual: Potato__healthy,
Predicted: Potato__Early_blight.
Confidence: 33.87%

Actual: Potato__Early_blight,
Predicted: Potato__Early_blight.
Confidence: 33.94%

Actual: Potato__Late_blight,
Predicted: Potato__Early_blight.
Confidence: 34.08%

Actual: Potato__Early_blight,
Predicted: Potato__Early_blight.
Confidence: 33.95%

Actual: Potato__Late_blight,
Predicted: Potato__Early_blight.
Confidence: 33.84%
```